Dear Hiring Team,

My name is Filippos Dimitrios Ktistakis, and I am applying for the Data Scientist position (R8114) in Copenhagen. I recently completed my MSc in Computational Physics at the University of Copenhagen, where I worked with data analysis, optimization, and machine learning on real-world problems.

In my thesis, I developed a misalignment detection tool for the McStas neutron scattering simulation package using optimization and machine learning methods. This involved building and validating neural-network-based models, designing robust data pipelines, and iterating until the results were stable enough to guide practical decisions. As a Data Analyst at the University of the Aegean, I processed complex environmental and microbiological datasets, focusing on data quality, cleaning, and converting raw measurements into clear, decision-ready insights for a broader team.

Across my studies, jobs, and independent projects, I have accumulated several years of hands-on experience working with machine learning models, neural networks, and data-driven tools from idea to deployment-ready prototypes. I am comfortable deploying and operating models in real applications, monitoring their behaviour, and troubleshooting performance issues when needed. My technical strengths are in Python and common data/ML libraries (NumPy, Pandas, SciPy, Matplotlib, TensorFlow, Keras, PyTorch, scikit-learn), together with a solid background in statistics, optimization, and time-series analysis.

Beyond Python, I have experience in C/C++, MATLAB, and parallel programming with MPI/OpenMP, and I have also worked with Julia in the context of high-performance scientific and quantum-optics style simulations, where performance profiling and GPU-accelerated methods are important. I enjoy optimizing systems and processes, whether that means speeding up numerical code, cleaning up data flows, or simplifying a model so that it is more robust in production. I like integrating AI components into broader systems and have built several applications where AI or time-series models are part of larger web or mobile products.

A large part of my free time goes into following AI developments and experimenting with new tools, agentic workflows, and integration patterns. On my portfolio (6x7.gr), you can see projects such as AI assistants, an AI-powered job application helper, forecasting tools, and other applications where machine learning is embedded in a concrete product. One long-term goal is building a general AI assistant that can coordinate and automate tasks across a user's device with minimal friction, and I am actively experimenting with combining different components and services toward that.

I enjoy working in a team and consider myself a reliable and practical team player who communicates clearly and keeps things honest and straightforward. I am fluent in English, and while my Danish is not yet fluent, I am currently in module 3 and working consistently to improve so that I can contribute more smoothly in a Danish-speaking environment. Intrum's focus on operational AI, integrating models into collections and decision platforms, and continuously testing and improving new AI use cases fits very well with the kind of work I like to do.

I would be glad to discuss how my experience in data science, optimization, and AI-driven product development can support Intrum Denmark's AI backbone and decision systems. Thank you for your time and consideration.

Kind regards,
Filippos Dimitrios Ktistakis